

LSTM Model on Amazon Fine Food Reviews Dataset

Exercise :

1. Download Amazon Fine Food Reviews dataset from Kaggle. (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)
2. Get vocabulary for each word in corpus.
3. Also get the frequencies for each word and index them from most frequent to less frequent.
4. Now run LSTM Models on the dataset.
5. Also try 2-layers of LSTM.
6. Also use dropout and batch normalization and plot train-test error vs epochs for each model.
7. Write your observations in English as crisply and unambiguously as possible. Always quantify your results.

Information regarding data set :

1. **Title:** Amazon Fine Food Reviews Data
2. **Sources:** Stanford Network Analysis Project(SNAP)
3. **Relevant Information:** This dataset consists of reviews of fine foods from amazon. The data span a period of more than 10 years, including all ~568,454 reviews up to October 2012(Oct 1999 - Oct 2012). Reviews include product and user information, ratings, and a plain text review.
4. **Attribute Information:**
 - ProductId** - unique identifier for the product
 - UserId** - unique identifier for the user
 - ProfileName** - name of the user
 - HelpfulnessNumerator** - number of users who found the review helpful
 - HelpfulnessDenominator** - number of users who indicated whether they found the review helpful or not
 - Score** - rating between 1 and 5.(rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored)
 - Time** - timestamp for the review
 - Summary** - brief summary of the review
 - Text** - text of the review

Objective :

It is a 2-class classification task, where we have to analyze, transform and perform LSTM to find the polarity of the dataset.

```
In [1]: import warnings
from sklearn.exceptions import DataConversionWarning
warnings.filterwarnings(action='ignore', category=DataConversionWarning)

import sqlite3
import datetime as dt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
from itertools import islice

from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.preprocessing import sequence
from keras.initializers import he_normal
from keras.layers import BatchNormalization, Dense, Dropout, Flatten, LSTM
from keras.layers.embeddings import Embedding
from keras.regularizers import L1L2
from prettytable import PrettyTable
```

Using TensorFlow backend.

Load Dataset

```
In [2]: # This dataset is already gone through data deduplication and text preprocessing, so it is approx ~364
        K

        # For Data Cleaning Steps follow this Link -
        # ipython notebook - https://drive.google.com/open?id=1JXCva5vXdIPgHbfNdD9sgnySqELoVtpty
        # dataset - https://drive.google.com/open?id=1IoDoTT8TfDu53N6cyKg6xVCU-FDPHyIF

        # For Text Preporcessing Steps follow this Link -
        # ipython notebook - https://drive.google.com/open?id=18-AkTzzEhCwM_hfLIbDNBMAP-imX4k4i
        # dataset - https://drive.google.com/open?id=1SfDwwXFhDpjgtfIE50_E80S089xRc8Sa

        # Load dataset
        def load_review_dataset(do_not_sample=True, sample_count=1):
            # Create connection object to load sqlite dataset
            connection = sqlite3.connect('finalDataSet.sqlite')

            # Load data into pandas dataframe.
            reviews_df = pd.read_sql_query(""" SELECT * FROM Reviews """,connection)

            # Drop index column
            reviews_df = reviews_df.drop(columns=['index'])

            # Sample dataset
            if do_not_sample == False:
                reviews_df = reviews_df.sample(sample_count)

            # Convert timestamp to datetime.
            reviews_df['Time'] = reviews_df[['Time']].applymap(lambda x: dt.datetime.fromtimestamp(x))

            # Sort the data on the basis of time.
            reviews_df = reviews_df.sort_values(by=['Time'])

            return reviews_df

        # Load 'finalDataSet.sqlite' in panda's daraframe.
        reviews_df = load_review_dataset(do_not_sample = True,sample_count = 1)

        # Make CleanedText as a dataset for clustering
        cleaned_text = reviews_df['CleanedText'].values

        print("Dataset Shape : \n",cleaned_text.shape)

        reviews_df['Score'] = reviews_df['Score'].map(lambda x : 1 if x == 'positive' else 0)
        reviews_df.head(5)
```

Dataset Shape :
(351237,)

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
382	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	1
250	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2	1
383	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	0	1
269	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2	2	1
369	374343	B00004CI84	A1B2IZU1JLZA6	Wes	19	23	0

Lets calculate frequencies for each word and index them from most frequent to less frequent.

```
In [3]: all_words=[]
for sentence in cleaned_text:
    words = sentence.split()
    all_words += words

print("Shape of the data : ",cleaned_text.shape)
print("Number of sentences present in complete dataset : ",len(all_words))

counts = Counter(all_words)
print("Number of unique words present in whole corpus: ",len(counts.most_common()))
vocab_size = len(counts.most_common()) + 1
top_words_count = 5000
sorted_words = counts.most_common(top_words_count)

word_index_lookup = dict()
i = 1
for word,frequency in sorted_words:
    word_index_lookup[word] = i
    i += 1

print()
print("Top 25 words with their frequencies:")
print(counts.most_common(25))
print()
print("Top 25 words with their index:")
print(list(islice(word_index_lookup.items(), 25)))
```

Shape of the data : (351237,)
Number of sentences present in complete dataset : 12901678
Number of unique words present in whole corpus: 93072

Top 25 words with their frequencies:
[('like', 160957), ('tast', 153682), ('flavor', 122605), ('good', 120139), ('love', 111232), ('use', 10705), ('product', 110217), ('one', 108864), ('great', 104928), ('tri', 96815), ('tea', 89507), ('coffe', 88109), ('get', 80051), ('make', 79567), ('food', 69353), ('would', 67655), ('buy', 63884), ('time', 60622), ('realli', 57989), ('eat', 57247), ('amazon', 55670), ('order', 55535), ('dont', 53855), ('much', 53351), ('price', 53027)]

Top 25 words with their index:
[('like', 1), ('tast', 2), ('flavor', 3), ('good', 4), ('love', 5), ('use', 6), ('product', 7), ('one', 8), ('great', 9), ('tri', 10), ('tea', 11), ('coffe', 12), ('get', 13), ('make', 14), ('food', 15), ('would', 16), ('buy', 17), ('time', 18), ('realli', 19), ('eat', 20), ('amazon', 21), ('order', 22), ('dont', 23), ('much', 24), ('price', 25)]

Lets add new column to our existing review dataframe with the index value of the words, which are present in 'CleanedText' columns.

```
In [4]: def apply_text_index(row):
        holder = []
        for word in row['CleanedText'].split():
            if word in word_index_lookup:
                holder.append(word_index_lookup[word])
            else:
                holder.append(0)
        return holder

reviews_df['CleanedText_Index'] = reviews_df.apply(lambda row: apply_text_index(row),axis=1)
reviews_df.head(5)
```

Out[4]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
382	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	1
250	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2	1
383	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	0	1
269	374422	B00004CI84	A1048CYU0OV4O8	Judy L. Eans	2	2	1
369	374343	B00004CI84	A1B2IZU1JLZA6	Wes	19	23	0

Split dataset into 70 : 30 split.


```

In [8]: # Batch size
batch_size = 192

# Number of time whole data is trained
epochs = 10

# Embedding vector size
embedding_vecor_length = 32

# Bias regularizer value - we will use elasticnet
reg = L1L2(0.01, 0.01)

# Plot train and cross validation loss
def plot_train_cv_loss(trained_model, epochs, colors=['b']):
    fig, ax = plt.subplots(1,1)
    ax.set_xlabel('epoch')
    ax.set_ylabel('Categorical Crossentropy Loss')
    x_axis_values = list(range(1,epochs+1))

    validation_loss = trained_model.history['val_loss']
    train_loss = trained_model.history['loss']

    ax.plot(x_axis_values, validation_loss, 'b', label="Validation Loss")
    ax.plot(x_axis_values, train_loss, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()

```

```

In [9]: # Instantiate sequential model
model = Sequential()

# Add Embedding Layer
model.add(Embedding(vocab_size, embedding_vecor_length, input_length=max_review_length))

# Add batch normalization
model.add(BatchNormalization())

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer
model.add(LSTM(100))

# Add dropout
model.add(Dropout(0.20))

# Add Dense Layer
model.add(Dense(1, activation='sigmoid'))

# Summary of the model
print("Model Summary: \n")
model.summary()
print()
print()

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Run the model
trained_model = model.fit(x_train, np.array(y_train), batch_size = batch_size, epochs = epochs, verbose=1, validation_data=(x_test, y_test))

```

Model Summary:

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 500, 32)	2978336
batch_normalization_1 (Batch Normalization)	(None, 500, 32)	128
dropout_1 (Dropout)	(None, 500, 32)	0
lstm_1 (LSTM)	(None, 100)	53200
dropout_2 (Dropout)	(None, 100)	0
dense_1 (Dense)	(None, 1)	101
Total params: 3,031,765		
Trainable params: 3,031,701		
Non-trainable params: 64		

Train on 245865 samples, validate on 105372 samples

Epoch 1/10

245865/245865 [=====] - 1138s 5ms/step - loss: 0.2057 - acc: 0.9195 - val_loss: 0.1938 - val_acc: 0.9251

Epoch 2/10

245865/245865 [=====] - 1146s 5ms/step - loss: 0.1639 - acc: 0.9360 - val_loss: 0.1851 - val_acc: 0.9261

Epoch 3/10

245865/245865 [=====] - 1304s 5ms/step - loss: 0.1476 - acc: 0.9425 - val_loss: 0.1841 - val_acc: 0.9277

Epoch 4/10

245865/245865 [=====] - 1446s 6ms/step - loss: 0.1359 - acc: 0.9470 - val_loss: 0.1870 - val_acc: 0.9296

Epoch 5/10

245865/245865 [=====] - 1332s 5ms/step - loss: 0.1255 - acc: 0.9512 - val_loss: 0.1954 - val_acc: 0.9288

Epoch 6/10

245865/245865 [=====] - 1400s 6ms/step - loss: 0.1225 - acc: 0.9524 - val_loss: 0.1967 - val_acc: 0.9271

Epoch 7/10

245865/245865 [=====] - 1423s 6ms/step - loss: 0.1179 - acc: 0.9541 - val_loss: 0.2025 - val_acc: 0.9263

Epoch 8/10

245865/245865 [=====] - 1415s 6ms/step - loss: 0.1146 - acc: 0.9557 - val_loss: 0.2105 - val_acc: 0.9261

Epoch 9/10

245865/245865 [=====] - 1405s 6ms/step - loss: 0.1059 - acc: 0.9592 - val_loss: 0.2069 - val_acc: 0.9279

Epoch 10/10


245865/245865 [=====] - 1437s 6ms/step - loss: 0.1008 - acc: 0.9616 - val_loss: 0.2125 - val_acc: 0.9277

```
In [10]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test accuracy: {0:.2f}%'.format(score[1]*100))
```

Test accuracy: 92.77%

```
In [11]: print()
print()

# Plot train and cross validation error
plot_train_cv_loss(trained_model, epochs)
```



After 1st epoch we got 92.51% accuracy, and if we train further we starts to overfit, as validation error does not decreases.

Model 2 - With 2 - LSTM Layers

```
In [12]: %%time
# Instantiate sequential model
model = Sequential()

# Add Embedding Layer
model.add(Embedding(vocab_size, embedding_vecor_length, input_length=max_review_length))

# Add batch normalization
model.add(BatchNormalization())

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer 1
model.add(LSTM(100, return_sequences=True))

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer 2
model.add(LSTM(100))

# Add dropout
model.add(Dropout(0.20))

# Add Dense Layer
model.add(Dense(1, activation='sigmoid'))

# Summary of the model
print("Model Summary: \n")
model.summary()
print()
print()

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Run the model
trained_model = model.fit(x_train, np.array(y_train), batch_size = batch_size, epochs = epochs, verbose=1, validation_data=(x_test, y_test))
```


Model Summary:

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 500, 32)	2978336
batch_normalization_2 (Batch Normalization)	(None, 500, 32)	128
dropout_3 (Dropout)	(None, 500, 32)	0
lstm_2 (LSTM)	(None, 500, 100)	53200
dropout_4 (Dropout)	(None, 500, 100)	0
lstm_3 (LSTM)	(None, 100)	80400
dropout_5 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 1)	101
Total params: 3,112,165		
Trainable params: 3,112,101		
Non-trainable params: 64		

Train on 245865 samples, validate on 105372 samples


```
Epoch 1/10
245865/245865 [=====] - 2416s 10ms/step - loss: 0.2031 - acc: 0.9202 - val_loss: 0.1907 - val_acc: 0.9230
Epoch 2/10
245865/245865 [=====] - 2653s 11ms/step - loss: 0.1621 - acc: 0.9366 - val_loss: 0.1783 - val_acc: 0.9281
Epoch 3/10
245865/245865 [=====] - 2531s 10ms/step - loss: 0.1457 - acc: 0.9434 - val_loss: 0.1800 - val_acc: 0.9289
Epoch 4/10
245865/245865 [=====] - 2466s 10ms/step - loss: 0.1343 - acc: 0.9478 - val_loss: 0.1920 - val_acc: 0.9276
Epoch 5/10
245865/245865 [=====] - 2469s 10ms/step - loss: 0.1246 - acc: 0.9519 - val_loss: 0.1864 - val_acc: 0.9287
Epoch 6/10
245865/245865 [=====] - 2468s 10ms/step - loss: 0.1166 - acc: 0.9549 - val_loss: 0.1965 - val_acc: 0.9292
Epoch 7/10
245865/245865 [=====] - 2485s 10ms/step - loss: 0.1095 - acc: 0.9578 - val_loss: 0.2001 - val_acc: 0.9277
Epoch 8/10
245865/245865 [=====] - 2506s 10ms/step - loss: 0.1041 - acc: 0.9597 - val_loss: 0.2028 - val_acc: 0.9274
Epoch 9/10
245865/245865 [=====] - 2527s 10ms/step - loss: 0.0990 - acc: 0.9622 - val_loss: 0.1951 - val_acc: 0.9278
Epoch 10/10
245865/245865 [=====] - 2541s 10ms/step - loss: 0.0952 - acc: 0.9636 - val_loss: 0.2113 - val_acc: 0.9267
Wall time: 6h 57min 46s
```

```
In [13]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test accuracy: {:.2f}%'.format(score[1]*100))
```

Test accuracy: 92.67%

```
In [14]: print()
print()

# Plot train and cross validation error
plot_train_cv_loss(trained_model, epochs)
```



After 1st epoch we got 92.30% accuracy, and if we train further we starts to overfit, as validation error does not decreases.

Model 3 - With 5 - LSTM Layers

```
In [9]: # Instantiate sequential model
model = Sequential()

# Add Embedding Layer
model.add(Embedding(vocab_size, embedding_vector_length, input_length=max_review_length))

# Add batch normalization
model.add(BatchNormalization())

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer 1
model.add(LSTM(100,return_sequences=True,bias_regularizer=reg))

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer 2
model.add(LSTM(80,return_sequences=True,bias_regularizer=reg))

# Add dropout
model.add(Dropout(0.20))

# Add LSTM Layer 3
model.add(LSTM(60,return_sequences=True,bias_regularizer=reg))

# Add dropout
model.add(Dropout(0.30))

# Add LSTM Layer 4
model.add(LSTM(40,return_sequences=True,bias_regularizer=reg))

# Add batch normalization
model.add(BatchNormalization())

# Add dropout
model.add(Dropout(0.40))

# Add LSTM Layer 5
model.add(LSTM(20))

# Add dropout
model.add(Dropout(0.50))

# Add Dense Layer
model.add(Dense(1, activation='sigmoid'))

# Summary of the model
print("Model Summary: \n")
model.summary()
print()
print()

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Run the model
trained_model = model.fit(x_train, np.array(y_train), batch_size = batch_size, epochs = epochs, verbose=1, validation_data=(x_test, y_test))
```

Model Summary:

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 500, 32)	2978336
batch_normalization_1 (Batch Normalization)	(None, 500, 32)	128
dropout_1 (Dropout)	(None, 500, 32)	0
lstm_1 (LSTM)	(None, 500, 100)	53200
dropout_2 (Dropout)	(None, 500, 100)	0
lstm_2 (LSTM)	(None, 500, 80)	57920
dropout_3 (Dropout)	(None, 500, 80)	0
lstm_3 (LSTM)	(None, 500, 60)	33840
dropout_4 (Dropout)	(None, 500, 60)	0
lstm_4 (LSTM)	(None, 500, 40)	16160
batch_normalization_2 (Batch Normalization)	(None, 500, 40)	160
dropout_5 (Dropout)	(None, 500, 40)	0
lstm_5 (LSTM)	(None, 20)	4880
dropout_6 (Dropout)	(None, 20)	0
dense_1 (Dense)	(None, 1)	21
Total params: 3,144,645		
Trainable params: 3,144,501		
Non-trainable params: 144		

Train on 245865 samples, validate on 105372 samples

Epoch 1/10

245865/245865 [=====] - 4146s 17ms/step - loss: 2.3025 - acc: 0.9107 - val_loss: 0.2145 - val_acc: 0.9182

Epoch 2/10

245865/245865 [=====] - 3659s 15ms/step - loss: 0.1803 - acc: 0.9323 - val_loss: 0.1930 - val_acc: 0.9227

Epoch 3/10

245865/245865 [=====] - 3656s 15ms/step - loss: 0.1625 - acc: 0.9389 - val_loss: 0.1959 - val_acc: 0.9268

Epoch 4/10

245865/245865 [=====] - 3654s 15ms/step - loss: 0.1507 - acc: 0.9431 - val_loss: 0.1977 - val_acc: 0.9220

Epoch 5/10

245865/245865 [=====] - 3650s 15ms/step - loss: 0.1404 - acc: 0.9475 - val_loss: 0.1885 - val_acc: 0.9283

Epoch 6/10

245865/245865 [=====] - 3648s 15ms/step - loss: 0.1331 - acc: 0.9504 - val_loss: 0.2017 - val_acc: 0.9257

Epoch 7/10

245865/245865 [=====] - 3664s 15ms/step - loss: 0.1255 - acc: 0.9525 - val_loss: 0.1880 - val_acc: 0.9279

Epoch 8/10

245865/245865 [=====] - 3659s 15ms/step - loss: 0.1192 - acc: 0.9559 - val_loss: 0.2126 - val_acc: 0.9275

Epoch 9/10

245865/245865 [=====] - 3665s 15ms/step - loss: 0.1145 - acc: 0.9572 - val_loss: 0.2138 - val_acc: 0.9257

Epoch 10/10


245865/245865 [=====] - 3751s 15ms/step - loss: 0.1092 - acc: 0.9594 - val_loss: 0.2201 - val_acc: 0.9252

```
In [10]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test accuracy: {:.2f}%'.format(score[1]*100))
```

Test accuracy: 92.52%

```
In [11]: print()
print()

# Plot train and cross validation error
plot_train_cv_loss(trained_model, epochs)
```



After 7th epoch, gap between train error and test error increases drastically, so to avoid overfitting we should train our model upto 7th epoch.

We can see that, as we increase number of LSTM layers chance of overfitting the model reduces.

```
In [2]: # Pretty table instance
ptable = PrettyTable()
ptable.title = "Comparison between different LSTM Models"
ptable.field_names = ['Number of LSTM Layers', 'Epoch', 'Testing Accuracy', 'Does Overfit']
ptable.add_row(["1", "10", "90.77", "Yes"])
ptable.add_row(["2", "10", "92.67", "Yes"])
ptable.add_row(["5", "10", "92.52", "No - Optimal Epoch value is 7"])

# Print pretty table values
print(ptable)
```

Comparison between different LSTM Models			
Number of LSTM Layers	Epoch	Testing Accuracy	Does Overfit
1	10	90.77	Yes
2	10	92.67	Yes
5	10	92.52	No - Optimal Epoch value is 7

Observations :

1. Tried different LSTM architectures on Amazon Fine Food Review Dataset.
2. 'sigmoid' is used as an activation function to develop LSTM network.
3. 'Adam' is used as an optimizer to develop LSTM network.
4. Introduced batch normalization and dropout in between hidden layers.

Note:

To avoid overfitting we can try below following measures:

We can increase the number of epochs to some reasonable number like 100 - 300 ,

We can introduce 'recurrent_regularizer' on LSTM layer, for different values of L1 or L2 or elasticnet ,

We can also try 'kernel_regularizer' on LSTM layer, for different values of L1 or L2 or elasticnet ,

We can combine CNN + MaxPooling + LSTM and observe the any decrease in validation error.